

## Preprocessing

### Missing values

The original 'breast-cancer-wisconsin' dataset has missing values in some columns, which are marked by a '?' symbol. To deal with this, we replaced missing values with the **mode** of the corresponding column(feature).

### ID column in Cancer data

Since the IDs of patients seemed irrelevant to our classification task, we excluded the first column of data which has patient IDs for problem 1 and the (i) part of problem 2. However, this column is included in part(ii) of problem 2, as we need 10 columns for setting  $m=10$ .

### Image data

We normalized the image data in problem 3 from the range  $[0, 255]$  to the range  $[0, 1]$  to avoid large numbers while calculating distances squares.

## 1 Problem 1: Adaboost

### Summary of Methods

In this problem, we use Ada-boosting with 1-level decision trees for classification. We use the following steps to implement this classifier:

1. Initialize weight for all datapoints to  $1/N$  where  $N$  is the number of training samples.
2. Then we train a binary decision stump on this data using the information gain as criteria in the following way:
  - (a) For each attribute  $i$ , we get a list of all unique values  $\{v_{i,1}, v_{i,2}, \dots\}$ .
  - (b) We try each of these values  $v_{i,j}$  as a threshold to split the data into 2 groups.
  - (c) Calculate the entropy for these splits using the formula

$$H(Y|x_i) = p(x_i < v_{i,j})H(Y|x_i < v_{i,j}) + p(x_i \geq v_{i,j})H(Y|x_i \geq v_{i,j})$$

where,

$$H(Y|x_i < v_{i,j}) = -p(Y=0) \log p(Y=0) - p(Y=1) \log p(Y=1)$$

(These probabilities are for data with  $x_i < v_{i,j}$ )

- (d) Select the splitting attribute  $i$  with the highest information gain. Gain can be calculated as

$$\text{Gain}_i = H(Y) - H(Y|x_i)$$

For comparing different attributes we can simply use:

$$\text{Gain}_i = 1 - H(Y|x_i)$$

3. With this decision stump, we calculate the error rate on training data  $\epsilon$ . This can be used to compute the estimator weight  $\alpha_t$ :

$$\alpha_t = \frac{1}{2} \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

4. This classifier (decision stump) and corresponding  $\alpha_t$  are stored in a list  $\mathcal{L}$ .

5.  $\alpha_t$  is also used to update the weights for the training data according to equation:

$$w_{t+1} = w_t \exp(-\alpha_t y y_{pred})$$

Weights are normalized by dividing with the sum of all weights:

$$w_{t+1} = \frac{w_{t+1}}{\text{sum}(w_{t+1})}$$

6. For the next iteration, a new dataset is generated by using these new weights  $w_{t+1}$  as a probability distribution to sample from the original dataset.

7. Repeat steps 2-6 with the new dataset until the required number of weak learners (100) are reached.

To predict a new data-point using the learned model  $\mathcal{L}$ , we use each stump  $G_t$  in  $\mathcal{L}$  to make predictions  $G_t(x)$ . The final output labels are obtained by a weighted sum of all classifiers' predictions:

$$y_{pred} = \text{sign} \left[ \sum_{t=1}^T \alpha_t G_t(x) \right]$$

## Results

For this problem we used 50% training-testing split. Using **100** weak decision stumps as classifiers, we got a minimum **train error-rate of 0.3%** and **test error-rate of 4.6%**. The plot of 'Error rate' against the 'Number of weak classifiers' used is shown below:

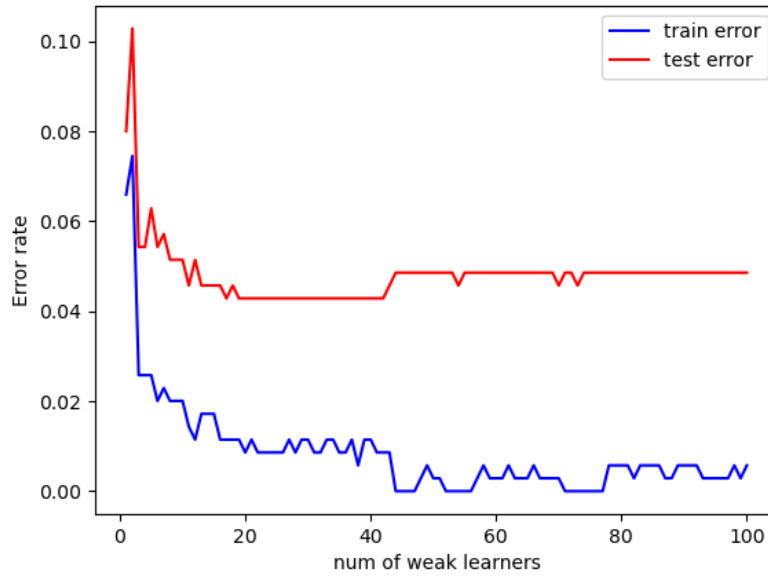


Figure 1: Adaboost: Error rate vs Number of weak learners

## 2 Problem 2: Random Forest Classification

In this problem we use a Random forest of 1-level decision stumps for classification. We use the following method for classification:

1. First we generate an equally sized random sample  $D_i$  from the given dataset  $D$ . (There is a chance some points will be repeated and some points will be omitted). We will use this sampled dataset  $D_i$  for the rest of the operations.
2. Out of all features in the data, we choose  $m = 3$  features, and train a decision stump by selecting a feature with highest information gain from just those  $m$  features.

3. Equations for information gain and entropy while selecting the splitting feature, are exactly the same as in the previous problem.
4. We store this decision stump classifier in a list  $\mathcal{L}$ .
5. Repeat the steps 1-4 for the desired number of trees (stumps) are learned.

To classify a new data-point  $x_i$ , we loop through all the decision stumps  $G_t$  stored in  $\mathcal{L}$ , and make predictions:  $G_t(x_i)$ . The final class label is given by the majority vote among all  $G_t(x_i)$ . Since the class labels are +1 and -1, we can just add all the predictions to see which prediction is more frequent:

$$y_{i,pred} = \text{sign} \left[ \sum_{t=1}^T G_t(x_i) \right]$$

## Results

First we run the random forest classifier with  $m = 3$  and increasing the number of trees from 1 to 100. Plot of error rate vs the number of trees is shown below:

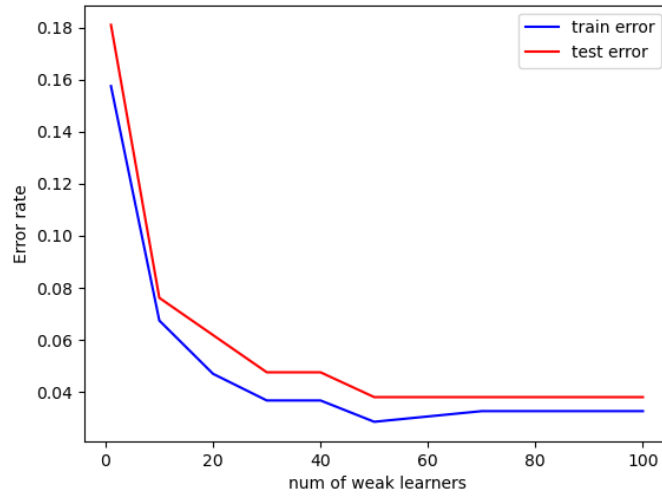


Figure 2: Random Forest: Error rate vs Number of trees

Next, we fix the number of trees as 100 and vary the number of attributes  $m$  selected for decision. Plot of error rate vs  $m$  is shown below:

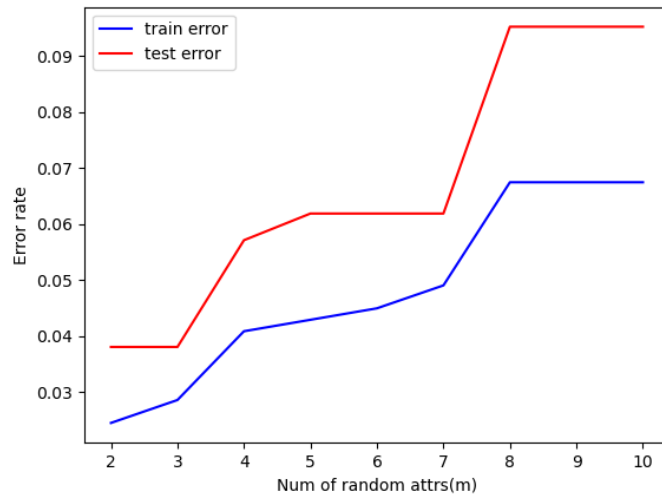


Figure 3: Random Forest: Error rate vs Number of randomly selected attributes

### 3 Problem 3: Image Segmentation using K-Means

In this problem, we use K-means clustering to reduce the number of colors used in a given image.

#### Summary of methods

1. Read image in as a  $(W \times H \times 3)$  numpy array. Reshape this into a  $(W.H \times 3)$  array to get a dataset in the shape of  $N \times D$ .
2. Initialize K random points  $\{C_1, C_2, C_3 \dots\}$  as cluster centers.
3. **Maximization step:** For each data point (pixel)  $x_n$ , calculate the distances from all K cluster centers. Assign the point to the cluster with the smallest distance.

$$r_{n,k} = \begin{cases} 1 & \text{if } k = \operatorname{argmin}_k (||x_n - C_k||^2) \\ 0 & \text{otherwise} \end{cases}$$

4. **Expectation step:** With the new cluster assignments, re-calculate the mean for each cluster.

$$C_k = \frac{\sum_{n=1}^N r_{n,k} ||x_n - C_k||^2}{\sum_{n=1}^N r_{n,k}}$$

If we find any cluster which does not have any points assigned to it, we re-initialize this cluster mean to a random (3D) point.

5. Calculate the overall distortion measure for all points  $J$ . It is the same as sum of squared distances of all points from their cluster centers.

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{n,k} ||x_n - C_k||^2$$

6. We say the convergence criteria is reached when the difference in this loss value between two iterations changes by less than **1%**. Until this convergence criteria is reached, repeat steps 3-4.

After convergence, we apply the final cluster centers' colors to all the points in that respective cluster. This will give us our final images.

#### Results

We run our K-means algorithm thrice using  $K= 3,5$  and  $7$ . We plot the training loss vs number of iterations for each run as shown below:





Figure 4: K-Means: Distortion Loss vs Number of Iterations

The resulting images after applying the respective cluster colors are shown in the next page:


(Please see next page)




(a) Original Image



(b)  $K = 3$



(c)  $K = 5$



(d)  $K = 7$

Figure 5: K-Means: Segmented/Compressed images